

Last Name:

First Name:

HW3 - Broadcast, Reduce and Allreduce

Rubric	Maximum Points	Points Received
Tree-based Communication Diagrams	20	
Tree-based fan-out broadcast and fan-in reduce	40	
Allreduce by fan-out broadcast, fan-in reduce	5	
All-to-one reduce and an one-to-all broadcast	5	
All programs working, complete report and answers to discussion questions	30	

Total: / 100

Please review the report for comments and grading feedback.

Additional grader remarks:

```

/*
 * Brandon Bell
 * csci4576
 * hw3 9-7-2016
 * Homework 3: Broadcast, Reduce, AllReduce.
 */

#include <stdio.h>
#include <string.h>
#include "mpi.h"

int MyPI_Bcast (
    void*      message,
    int*       count,
    MPI_Datatype datatype,
    int        root,
    MPI_Comm   comm,
    int        bitOrder
)

{
    // Determine the number of stages required for the tree.
    // I'm nabbing the Ceiling_log2() algorithm from chapter 5 in the book to
    // handle this.
    int p;
    int my_rank;
    MPI_Status status;
    MPI_Comm_size ( comm, &p );
    MPI_Comm_rank ( MPI_COMM_WORLD, &my_rank );
    unsigned int stages = 0;
    unsigned int temp = (unsigned) p - 1;
    unsigned int companion;

    while ( temp != 0 )
    {
        temp = temp >> 1;
        stages = stages + 1;
    }

    // Re-map mpi rank into a rankspace where rank 0 is the root. The broadcast
    // algorithm is based on bit arithmetic for a root with value zero and this
    // process zero must have the data beinf broadcast so, re-map the mpi rank
    // to root = rank zero with modulo arithmetic, the ultimate mapping is
    // irrelevant ( I think) as long as it is consistent and the new rank zero
    // has the message to be sent.
    int rank = ( my_rank + ( p - root ) ) % p;
    int mpirank; // = ( rank + ( p + root ) ) % p;
    int mpic;

    // Set up the loop to determine send and recvs for each stage of the tree.
    // This is for low to high bit method.
    for ( int stage = 0; stage < stages; stage++ )
    {
        // Determine the current processes send/Recv companion.
        companion = ( 1 << stage ) ^ rank;

        // Detemrine wheither to send/recv or do nothing with it's companion.
        // First determine if p is to recv at this stage.
        if ( ( rank < companion ) && ( companion >> ( 1 << stage ) ) == 0 )
        {
            mpic = ( companion + ( p + root ) ) % p;
            mpirank = ( rank + ( p + root ) ) % p;
        }
    }
}

```

```

printf(" send mpirank %d mpic %d\n", mpirank, mpic );
MPI_Send( message, count, MPI_FLOAT, mpirank, 1, comm );
printf("Step %d rank %d Sends to %d\n", stage, rank, companion);
}
else if ( ( rank > companion ) && ( companion >> ( 1 << stage ) ) == 0 )
{
    mpic = ( companion + ( p + root ) ) % p;
    mpirank = ( rank + ( p + root ) ) % p;
    printf(" recv mpirank %d mpic %d\n", mpirank, mpic );
    MPI_Recv( message, count, MPI_FLOAT, mpirank, 1, comm, &status );
    printf("Step %d rank %d Recvs from %d\n", stage, rank, companion);
}
// Or if p is supposed to send at this stage.
// Handle the case that p and it's companion are not to communicate
// during this stage.
else
{
    mpirank = ( companion + ( p + root ) ) % p;
    mpirank = ( rank + ( p + root ) ) % p;
    printf("Step %d rank %d does nothing\n", stage, mpirank);
    continue;
}
}
return message;
}

void MyPI_Reduce (
    void*      operand,
    void*      result,
    int        count,
    MPI_Datatype datatype,
    MPI_Op     operator,
    int        root,
    MPI_Comm   comm,
    int        bitOrder
)
{
}

void MyPI_AllReduce (
    void*      operand,
    void*      result,
    int        count,
    MPI_Datatype datatype,
    MPI_Op     operator,
    int        root,
    MPI_Comm   comm,
    int        bitOrder
)
{
}

void MyPI_AllReduce_Trivial (
    void*      operand,
    void*      result,
    int        count,
    MPI_Datatype datatype,
    MPI_Op     operator,
    int        root,
    MPI_Comm   comm,
    int        bitOrder
)
{
}

```

BellBrandon_HW3.c

```
    }
}

int main ( int argc, char* argv[] )
{
    int    my_rank;
    int    p;
    int    source;
    int    dest;
    int    tag = 0;
    float  message;
    MPI_Status  status;

    // Spin-up Mpi.
    MPI_Init (&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &p);

    // Check to make sure that more than process are running.
    if ( p == 1 )
    {
        printf("Please Run with > 1 processes.\n" );
        return 1;
    }

    if ( my_rank == 0 )
        message = 5;
    else
        message = 0;

    //printf("Message on rank %d is %f \n", my_rank, message );

    MPI_Bcast ( &message, 1, MPI_FLOAT, 0, MPI_COMM_WORLD, 1 );

    printf("Message on rank %d is %f \n", my_rank, message );

    MPI_Finalize();
    return 0;
}
```