

Last Name:

First Name:

HW4 - Butterfly Allreduce and Performance Analysis

Rubric	Maximum Points	Points Received
Tree-based “butterfly” Allreduce Diagrams	10	
Tree-based “butterfly” Allreduce Routines	20	
Instrumentation and Timing	30	
Plots and Performance Analysis	30	
Report and Discussion Questions	10	

Total: / 100

Please review the report for comments and grading feedback.

Additional grader remarks:

```

/*
 * Brandon Bell
 * csci4576
 * hw4 9-14-2016
 * Homework 4: Butterfly AllReduce and Performance Analysis.
 */

#include <stdio.h>
#include <string.h>
#include "mpi.h"

int MyMPI_Bcast ( void*
int*
MPI_Datatype
int
MPI_Comm
int
)
{
    message,
    count,
    MPI_Datatype
    datatype,
    int
    root,
    MPI_Comm
    comm,
    bitOrder
    bitOrder
}

// Determine the number of stages required for the tree
// I'm nabbing the Ceiling_log2() algorithm from chapter 5 in the book to
// handle this.
int p;
int my_rank;
MPI_Status
MPI_Comm_size ( comm, &p );
MPI_Comm_rank ( MPI_COMM_WORLD, &my_rank );
unsigned int stages = 0;
unsigned int temp = (unsigned) p - 1;
unsigned int companion;

while ( temp != 0 )
{
    temp = temp >> 1;
    stages = stages + 1;
}

// Re-map mpi rank into a rankspace where rank 0 is the root. The broadcast
// algorithm is based on bit arithmetic for a root with value zero and this
// process zero must have the data bein broadcast so, re-map the mpi rank
// to root = rank zero with modulo arithmetic, the ultimate mapping is
// irrelevant ( I think) as long as it is consistent and the new rank zero
// has the message to be sent.
int rank = ( my_rank + ( p - root ) ) % p;
int mpirank; // = ( rank + ( p + root ) ) % p;
int mpic;

// Set up the loop to determine send and recvs for each stage of the tree.
// This is for low to high bit method.
for ( int stage = 0; stage < stages; stage++ )
{
    // Determine the current processes send/Recv companion.
    companion = ( 1 << stage ) ^ rank;

    // Detemine wheither to send/recv or do nothing with it's companion.
    // First determine if p is to recv at this stage.
    if ( ( rank < companion ) && ( companion >> ( 1 << stage ) ) == 0 )
    {
        mpic = ( companion + ( p + root ) ) % p;
        mpirank = ( rank + ( p + root ) ) % p;
    }
}

```

```

printf( " send mpirank %d mpic %d\n", mpirank, mpic );
MPI_Send( message, count, MPI_FLOAT, mpirank, 1, comm );
printf( "Step %d rank %d Sends to %d\n", stage, rank, companion );
}
else if ( ( rank > companion ) && ( companion >> ( 1 << stage ) ) == 0 )
{
    mpic = ( companion + ( p + root ) ) % p;
    mpirank = ( rank + ( p + root ) ) % p;
    printf( " recv mpirank %d mpic %d\n", mpirank, mpic );
    MPI_Recv( message, count, MPI_FLOAT, mpirank, 1, comm, &status );
    printf( "Step %d rank %d Recvs from %d\n", stage, rank, companion );
}
// Or if p is supposed to send at this stage.
// Handle the case that P and it's companion are not to communicate
// during this stage.
else
{
    mpirank = ( companion + ( p + root ) ) % p;
    mpirank = ( rank + ( p + root ) ) % p;
    printf( "Step %d rank %d does nothing\n", stage, mpirank );
    continue;
}
}
return message;
}

void MyMPI_Reduce ( void*
void*
int
MPI_Datatype
MPI_Op
operator,
int
root,
MPI_Comm
comm,
bitOrder
bitOrder
)
{
}

void MyMPI_AllReduce ( void*
void*
int
MPI_Datatype
MPI_Op
operator,
int
root,
MPI_Comm
comm,
bitOrder
bitOrder
)
{
}

void MyMPI_AllReduce_Trivial ( void*
void*
int
MPI_Datatype
MPI_Op
operator,
int
root,
MPI_Comm
comm,
bitOrder
bitOrder
)
{
}

```

```
{
    MPI_Status status;
    MPI_Comm_size ( comm, &p );
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    // P0 only recvs.
    if (my_rank == 0)
    {
        source = 1;
        MPI_Recv( &result, count, MPI_FLOAT, source, tag, comm, &status );
        total = total + integral;
    }
    //
    else
    {
        source = 0;
        MPI_Send( &result, count, MPI_FLOAT, dest, tag, MPI_COMM_WORLD );
    }
}

int main ( int argc, char* argv[] )
{
    int my_rank;
    int p;
    int source;
    int dest;
    int tag = 0;
    float message;
    MPI_Status status;

    // Spin-up MPI.
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
    MPI_Comm_size(MPI_COMM_WORLD, &p);

    // Check to make sure that more than process are running.
    if ( p == 1 )
    {
        printf("Please Run with > 1 processes.\n" );
        return 1;
    }

    if ( my_rank == 0 )
        message = 5;
    else
        message = 0;

    //printf("Message on rank %d is %f \n", my_rank, message );

    MPI_Bcast ( &message, 1, MPI_FLOAT, 0, MPI_COMM_WORLD, 1 );

    printf("Message on rank %d is %f \n", my_rank, message );

    MPI_Finalize();
    return 0;
}
```